

14.332.435/16.332.530
Introduction to Deep Learning

Lecture 2
Machine Learning Basics (1)

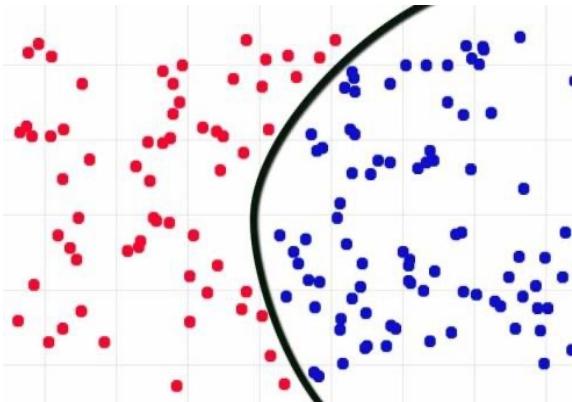
Yuqian Zhang
Department of Electrical and Computer Engineering

What does Learning Mean

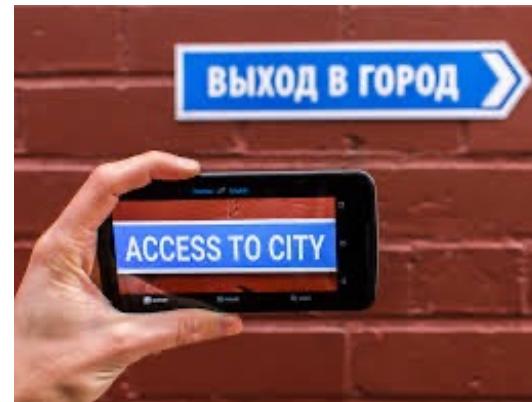
- ML algorithm is an algorithm able to *learn* from data
- A computer program is said to learn if
 - its **performance P** on **task T** can improve with **experience E**
- **Task:** What is ML algorithm used for
- **Performance:** Quantitative metric to evaluate algorithm
- **Experience:** Effect of ML algorithm on dataset

Example Tasks of Machine Learning

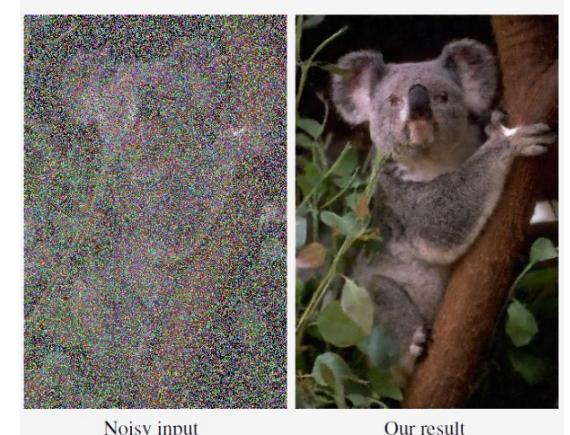
Classification



Translation



Denoise



Performance Metrics of ML

- Different tasks T have different performance measure P

Image Classification

TABLE II: AlexNet on ImageNet [4]. PD: Permuted Diagonal.

AlexNet	Block size (p) for PD weight matrix of FC6-FC7-FC8	Top-5 Acc.	Compression for overall FC layers
Original 32-bit float	1-1-1	80.20%	234.5MB(1 \times)
32-bit float with PD	10-10-4	80.00%	25.9MB(9.0 \times)
16-bit fixed with PD	10-10-4	79.90%	12.9MB(18.1 \times)

Language Translation

TABLE III: Stanford NMT (32-FC layer LSTMs) on IWSLT15 for English-Vietnamese Translation.

Stanford NMT (32-FC layer LSTMs) ⁵	Block size (p) for PD weight matrix of ALL FC Layers	BLEU Points	Compression for overall FC layers
Original 32-bit float	1	23.3	419.4MB(1 \times)
32-bit float with PD	8	23.3	52.4MB(8 \times)
16-bit fixed with PD	8	23.2	26.2MB(16 \times)

Taxonomy of Machine Learning: Different Experience

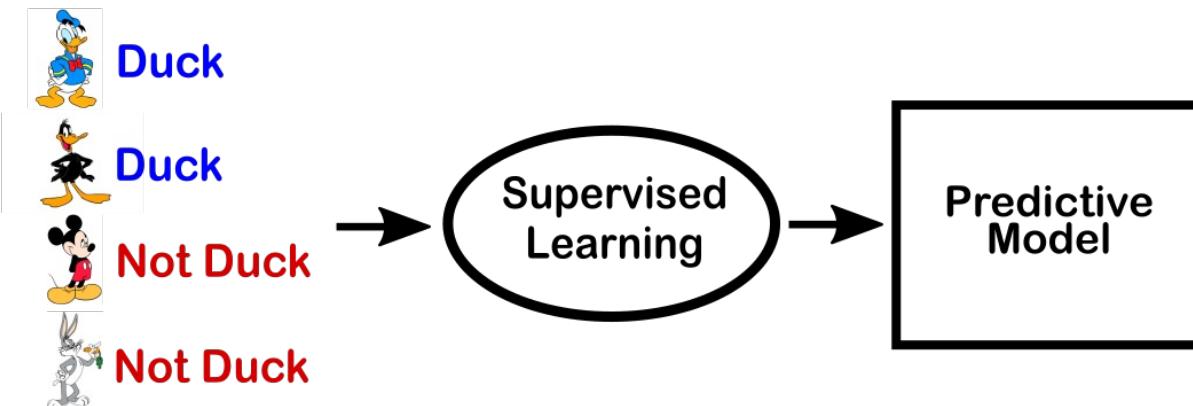
According to what kind of experience they are allowed to have during the learning process, ML algorithms can be categorized as:

- Supervised learning
- Unsupervised learning
- Reinforcement learning

Supervised learning

Supervised learning

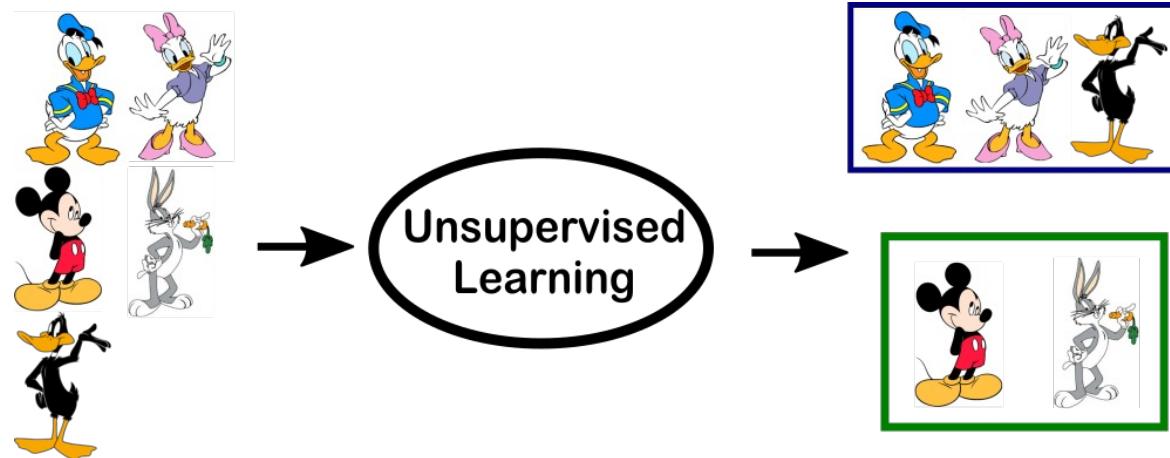
- Learn from labeled examples
- Each example is a pair consisting of an input object and a desired output value (a.k.a supervisory signal).



Unsupervised learning

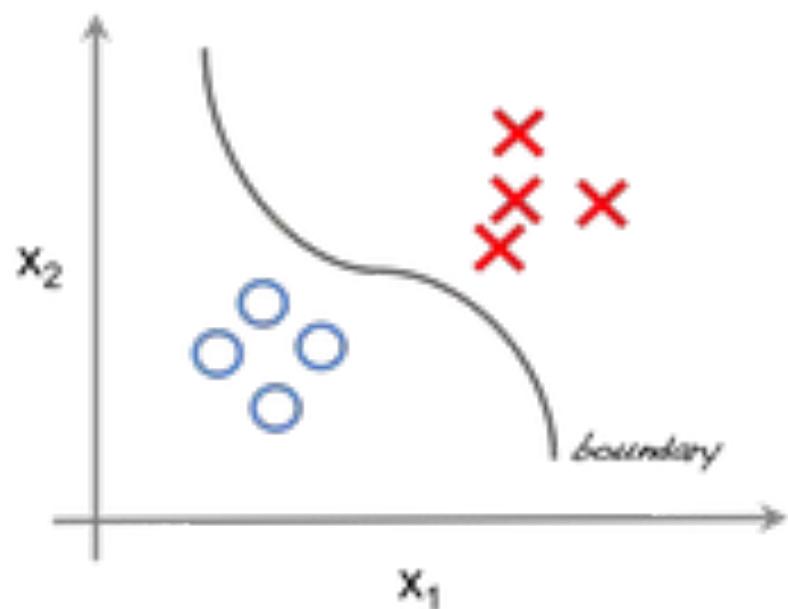
Unsupervised learning

- Draw inferences from datasets consisting of input data **without** labeled responses
- Find hidden patterns or grouping in data

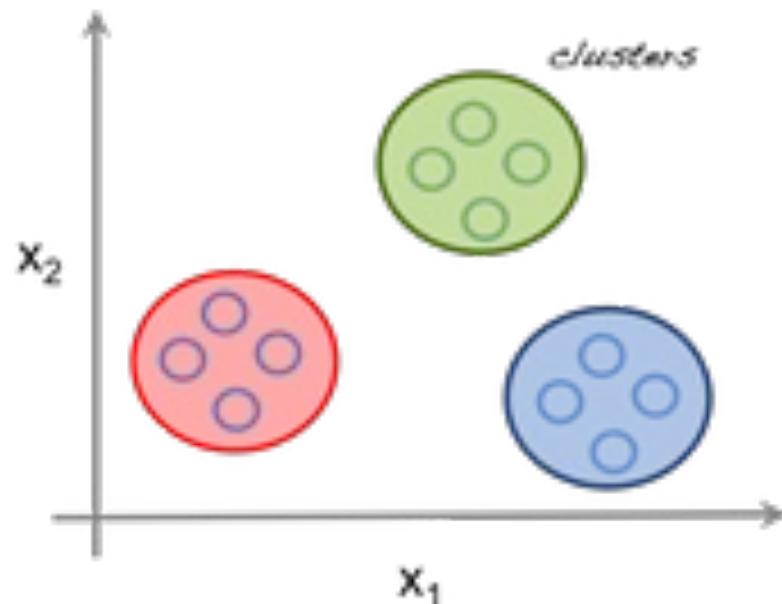


Supervised vs Unsupervised Learning

Supervised learning



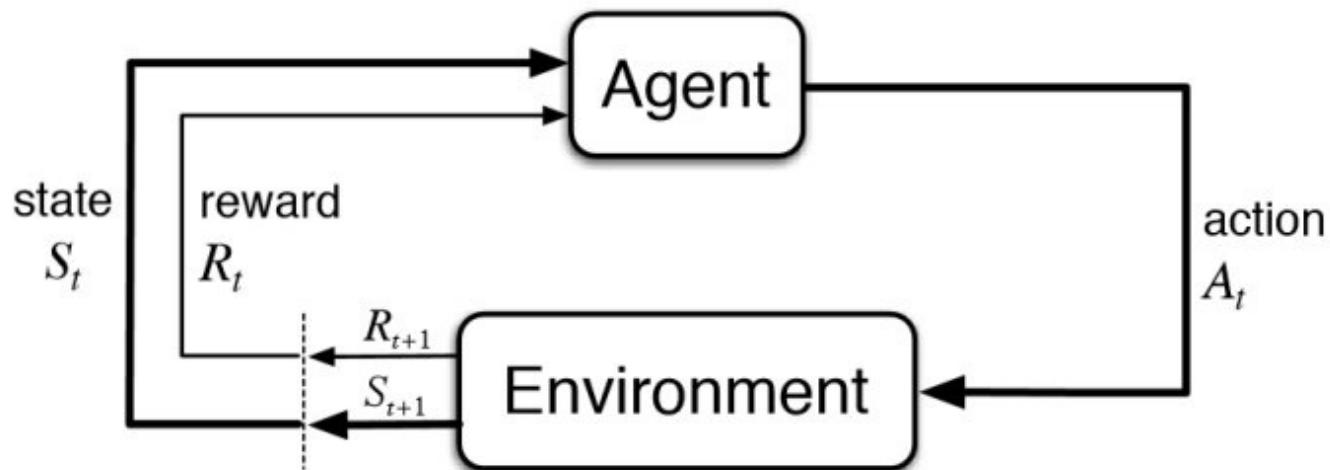
Unsupervised learning



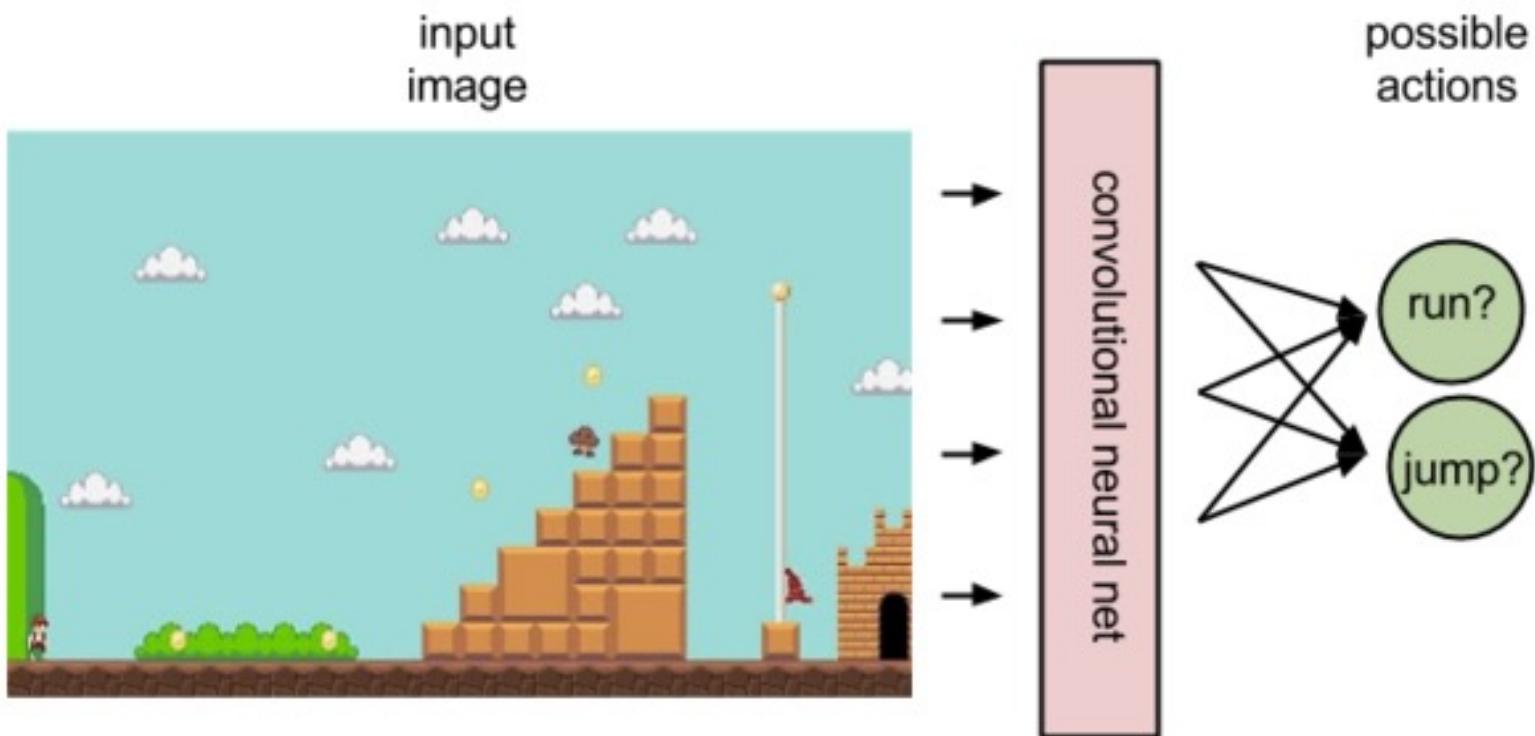
Reinforcement learning

Reinforcement learning

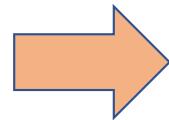
- Experience dataset and **interact with environment**
- **Feedback loop** between ML system and environment



Convolutional Agent

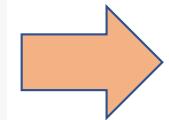


Recall Prior Example



An image classifier

```
def classify_image(image):  
    # Some magic here?  
    return class_label
```



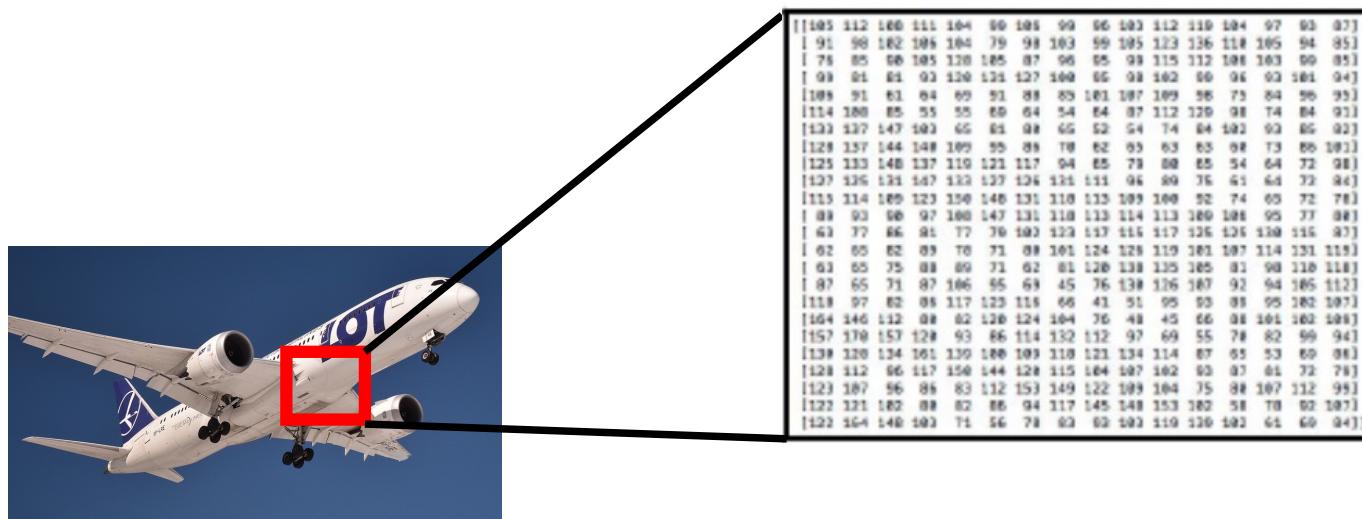
Plane

In this course we usually take image classification in computer vision as example

- Computer vision is the most successful field for DNNs
- Image classification is the core task of computer vision

Challenge: Semantic Gap

- In computer, each pixel is represented with a number
- Computer views an image as a big grid of numbers



Computer vision is not human vision

Challenge: Different Sizes, Shapes, Angles...



Machine Learning: Data-driven Approach

```
def train(images, labels):
    # Machine Learning!
    return model
```

```
def predict(model, test_images):
    # Use model to predict labels
    return test_labels
```

Example training set

airplane



automobile



bird



cat



deer



1. Collect a dataset of images and labels
2. Use Machine Learning to train a classifier
3. Evaluate the classifier on new images

A Naive Idea: Find Similar Images

```
def train(images, labels):
    # Machine Learning!
return model
```



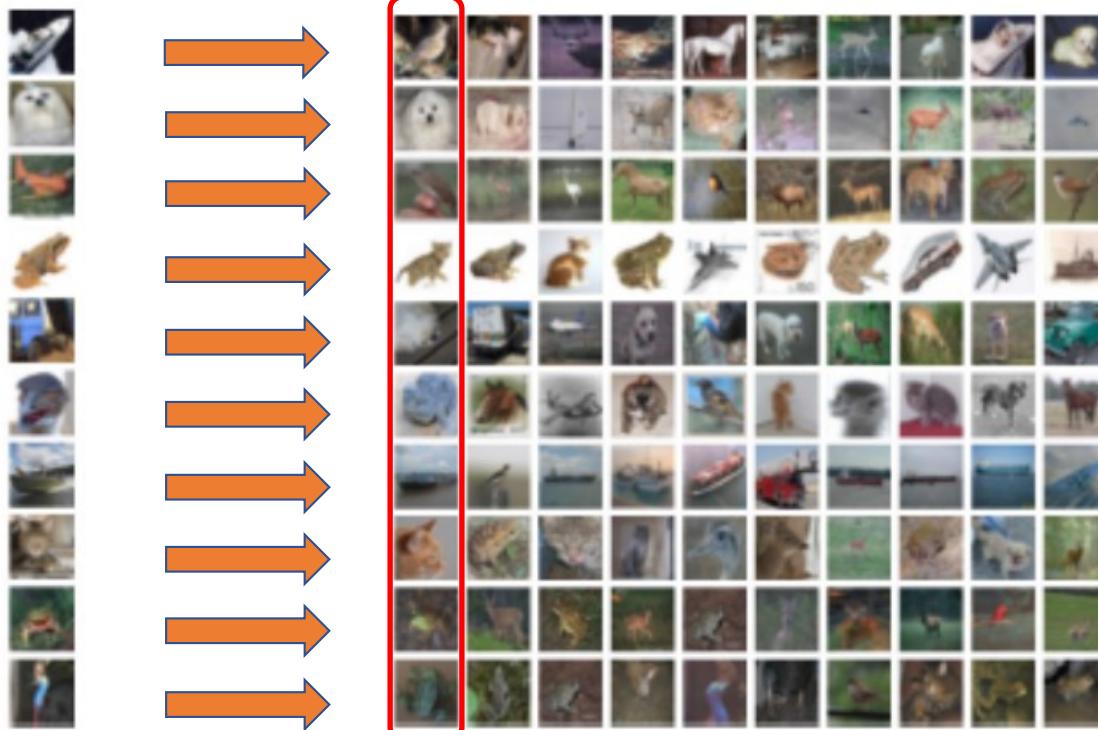
Memorize all data and labels

```
def predict(model, test_images):
    # Use model to predict labels
return test_labels
```



Predict the label of the most similar training images

Test image **Nearest Neighbors in training images**



How to Determine the Nearest?

L-distance is good metric:

e.g.: L1 distance: $d_1(I_1, I_2) = \sum_p |I_1^p - I_2^p|$ p : dimension

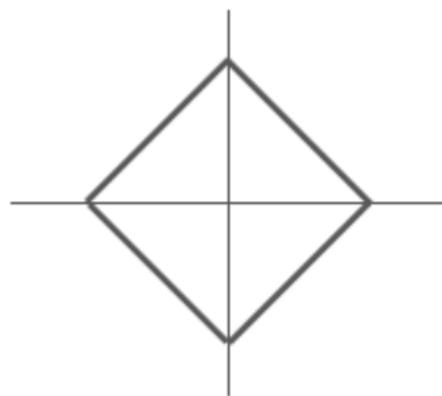
test image				training image				ρ pixel-wise absolute value differences			
56	32	10	18	10	20	24	17	46	12	14	1
90	23	128	133	8	10	89	100	82	13	39	33
24	26	178	200	12	16	178	170	12	10	0	30
2	0	255	220	4	32	233	112	2	32	22	108

add → 456

Distance Metric

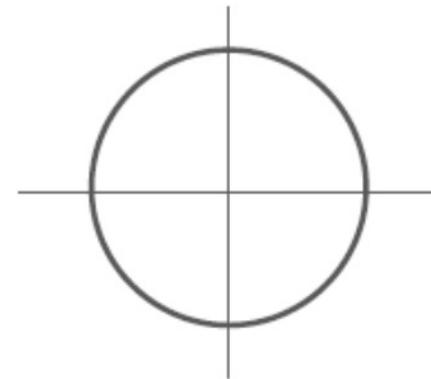
L1(Manhattan) distance

$$d_1(I_1, I_2) = \sum_p |I_1^p - I_2^p|$$



L2(Euclidean) distance

$$d_2(I_1, I_2) = \sqrt{\sum_p (I_1^p - I_2^p)^2}$$



Example Python Codes

```
import numpy as np

class NearestNeighbor:
    def __init__(self):
        pass

    def train(self, X, y):
        """ X is N x D where each row is an example. Y is 1-dimension of size N """
        # the nearest neighbor classifier simply remembers all the training data
        self.Xtr = X
        self.ytr = y

    def predict(self, X):
        """ X is N x D where each row is an example we wish to predict label for """
        num_test = X.shape[0]
        # lets make sure that the output type matches the input type
        Ypred = np.zeros(num_test, dtype = self.ytr.dtype)

        # loop over all test rows
        for i in xrange(num_test):
            # find the nearest training image to the i'th test image
            # using the L1 distance (sum of absolute value differences)
            distances = np.sum(np.abs(self.Xtr - X[i,:]), axis = 1)
            min_index = np.argmin(distances) # get the index with smallest distance
            Ypred[i] = self.ytr[min_index] # predict the label of the nearest example

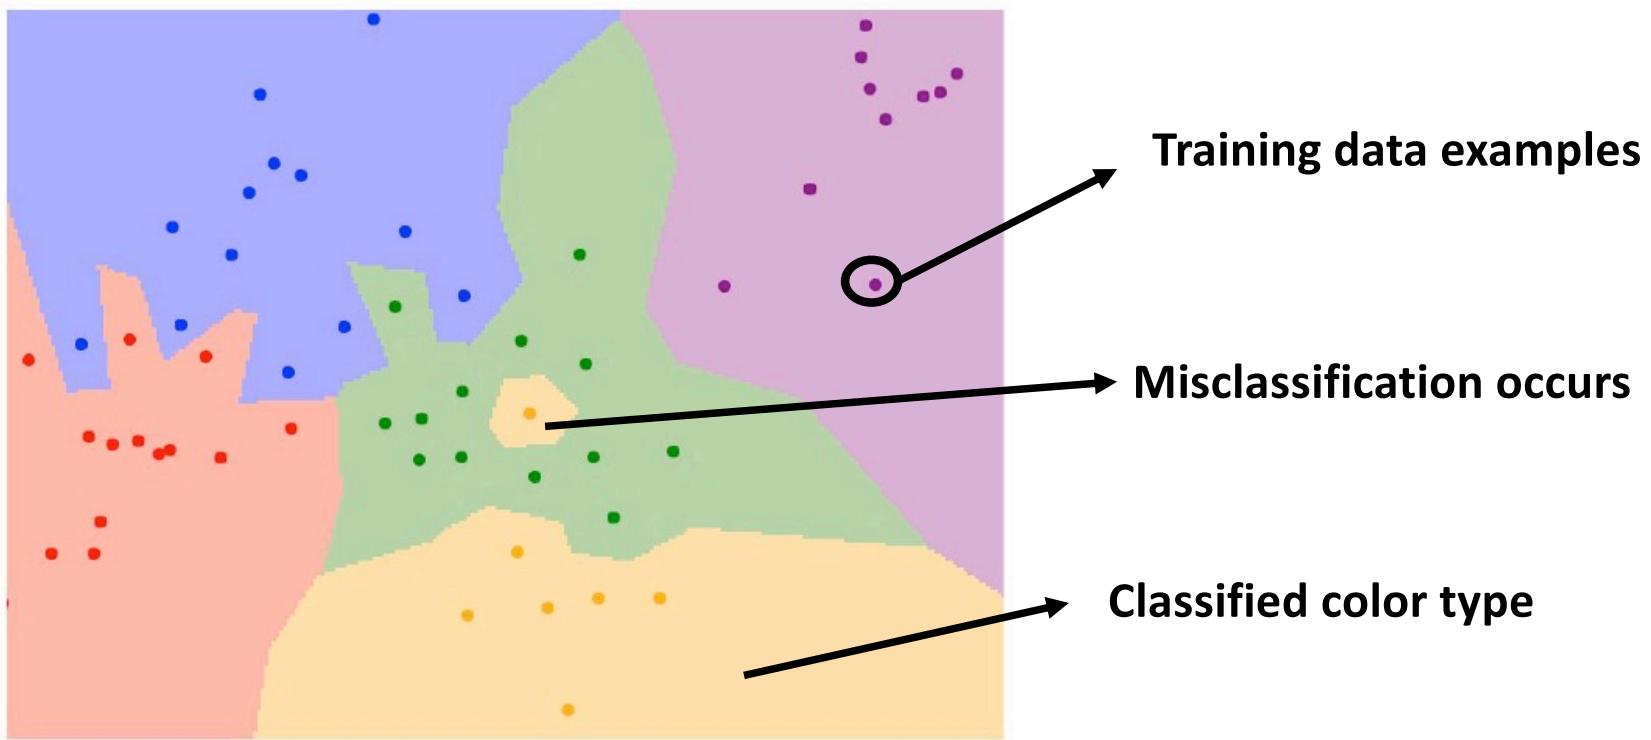
        return Ypred
```

With N training examples, time complexity for
Train is O(1) ----> Quick
Prediction is O(N) ----> Slow

Memorize training data

For each test image
Find nearest train image
Predict label of nearest image

Disadvantage If Only Considering Nearest

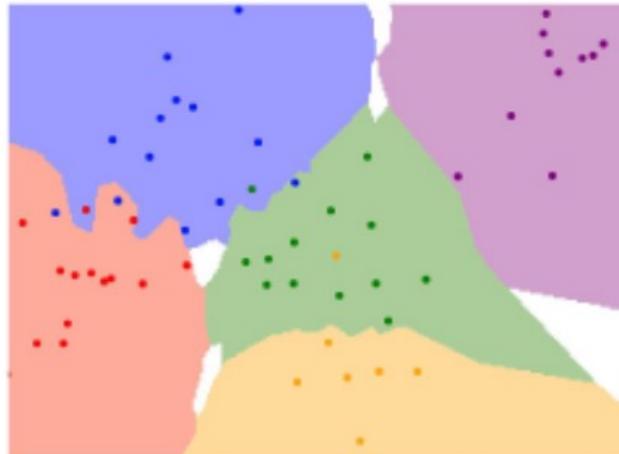


Improvement: Using K Nearest Neighbors (KNN)

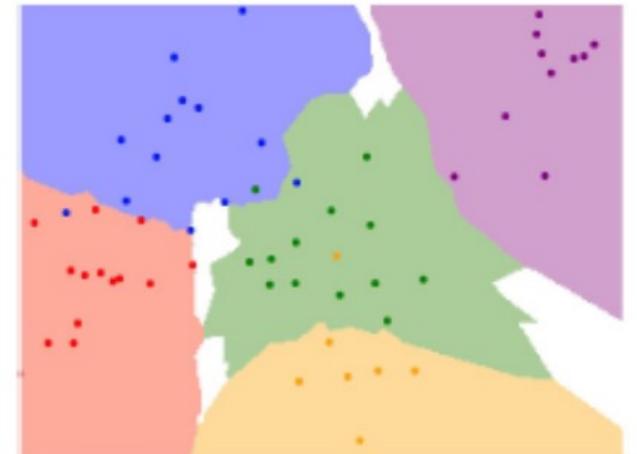
- KNN: Take **majority vote** from K closest points



$K = 1$



$K = 3$



$K = 5$

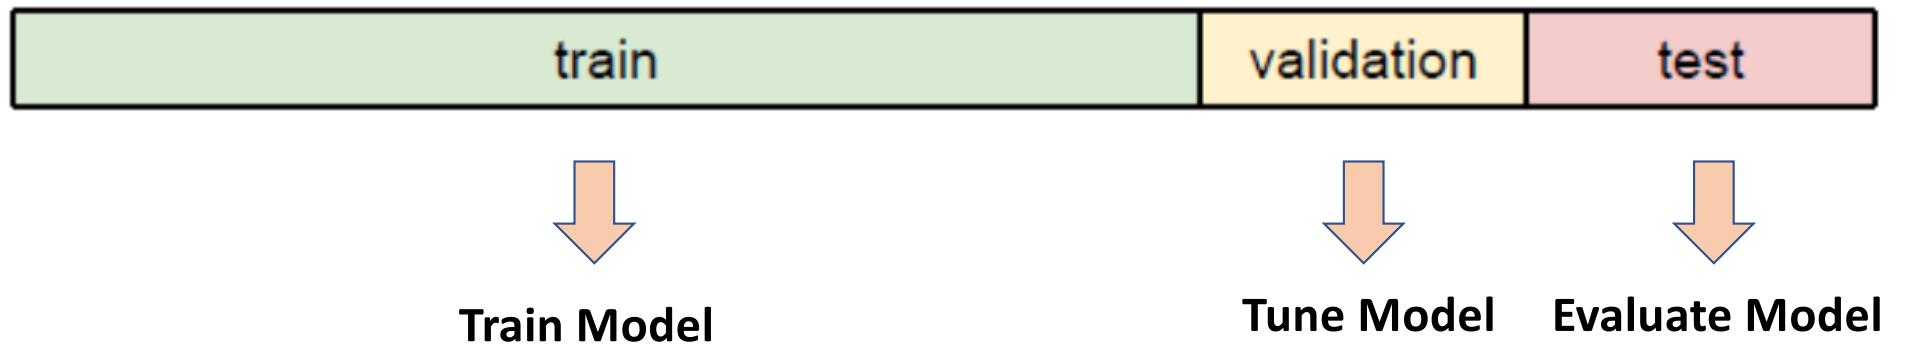
Hyperparameter

- When setting up KNN, you can choose two parameters:
 - What is the best value of k to use for voting?
 - What is the best distance to use for measuring?
- These are hyperparameters, which are not adapted by the machine learning algorithm itself.
- In DNN, many hyperparameters
 - # of layers/neuron, learning rate, weight decay.....

Train, Validation and Test Dataset

- Training set
 - a set of data used to discover predictive relationships
 - simple words: help to learn the model
- Test set
 - a set of data used to assess the strength and utility of a predictive relationship
 - simple words: help to evaluate a learned model
- Validation set
 - a set of examples used to tune the hyperparameters

Train, Validation and Test Dataset



Never use test data to train model!

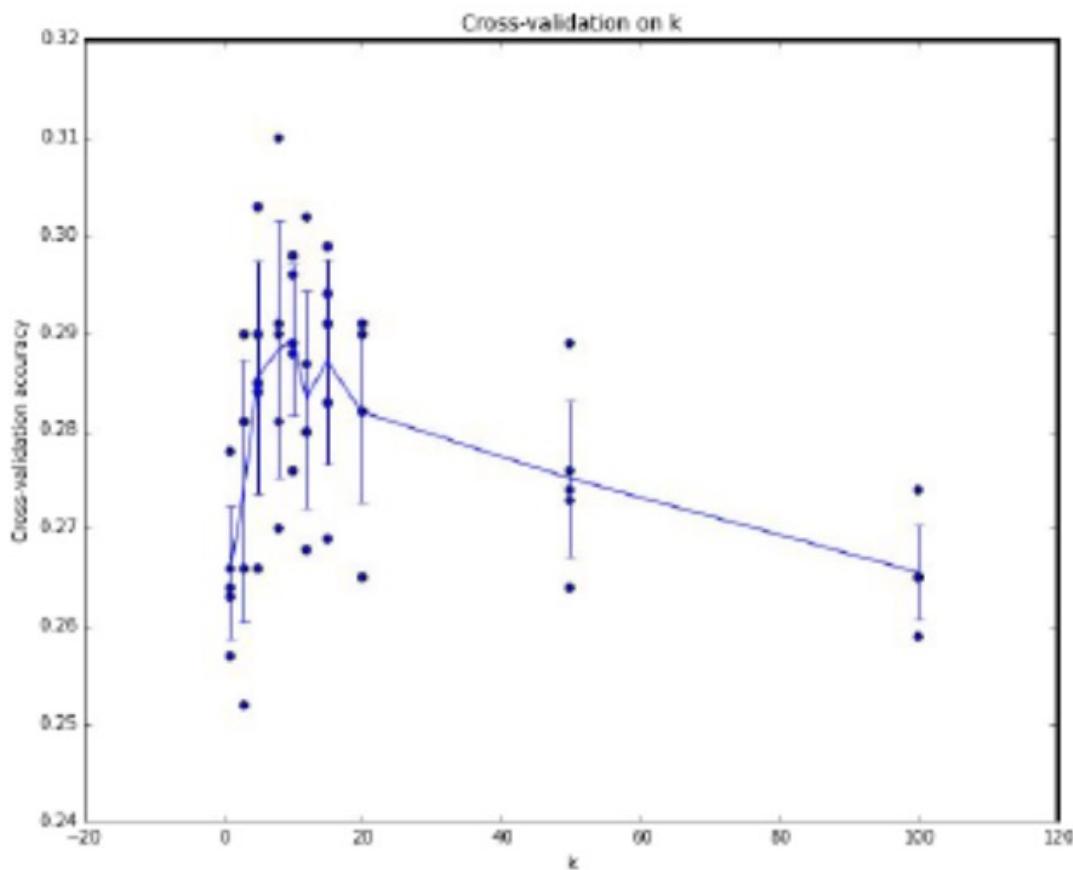
Cross-validation

When dataset is small

- Split data, try each fold as validation, and average
- Not too frequent in deep learning

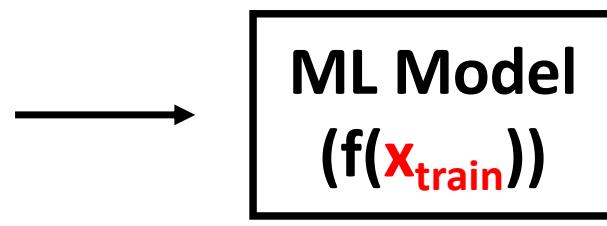
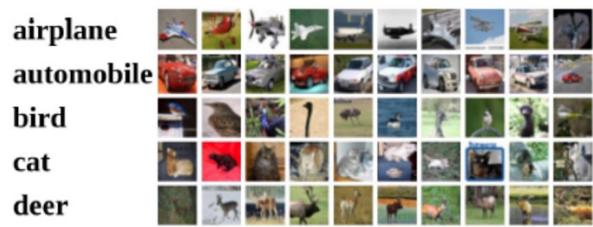
fold 1	fold 2	fold 3	fold 4	fold 5	test
fold 1	fold 2	fold 3	fold 4	fold 5	test
fold 1	fold 2	fold 3	fold 4	fold 5	test

5-fold Cross-validation for KNN



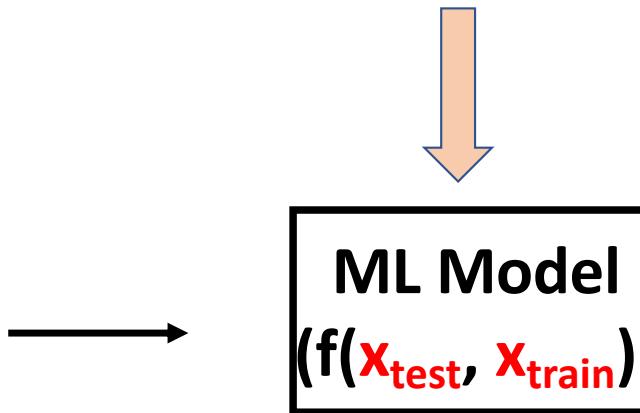
Summary of Non-parametric KNN

Training Examples



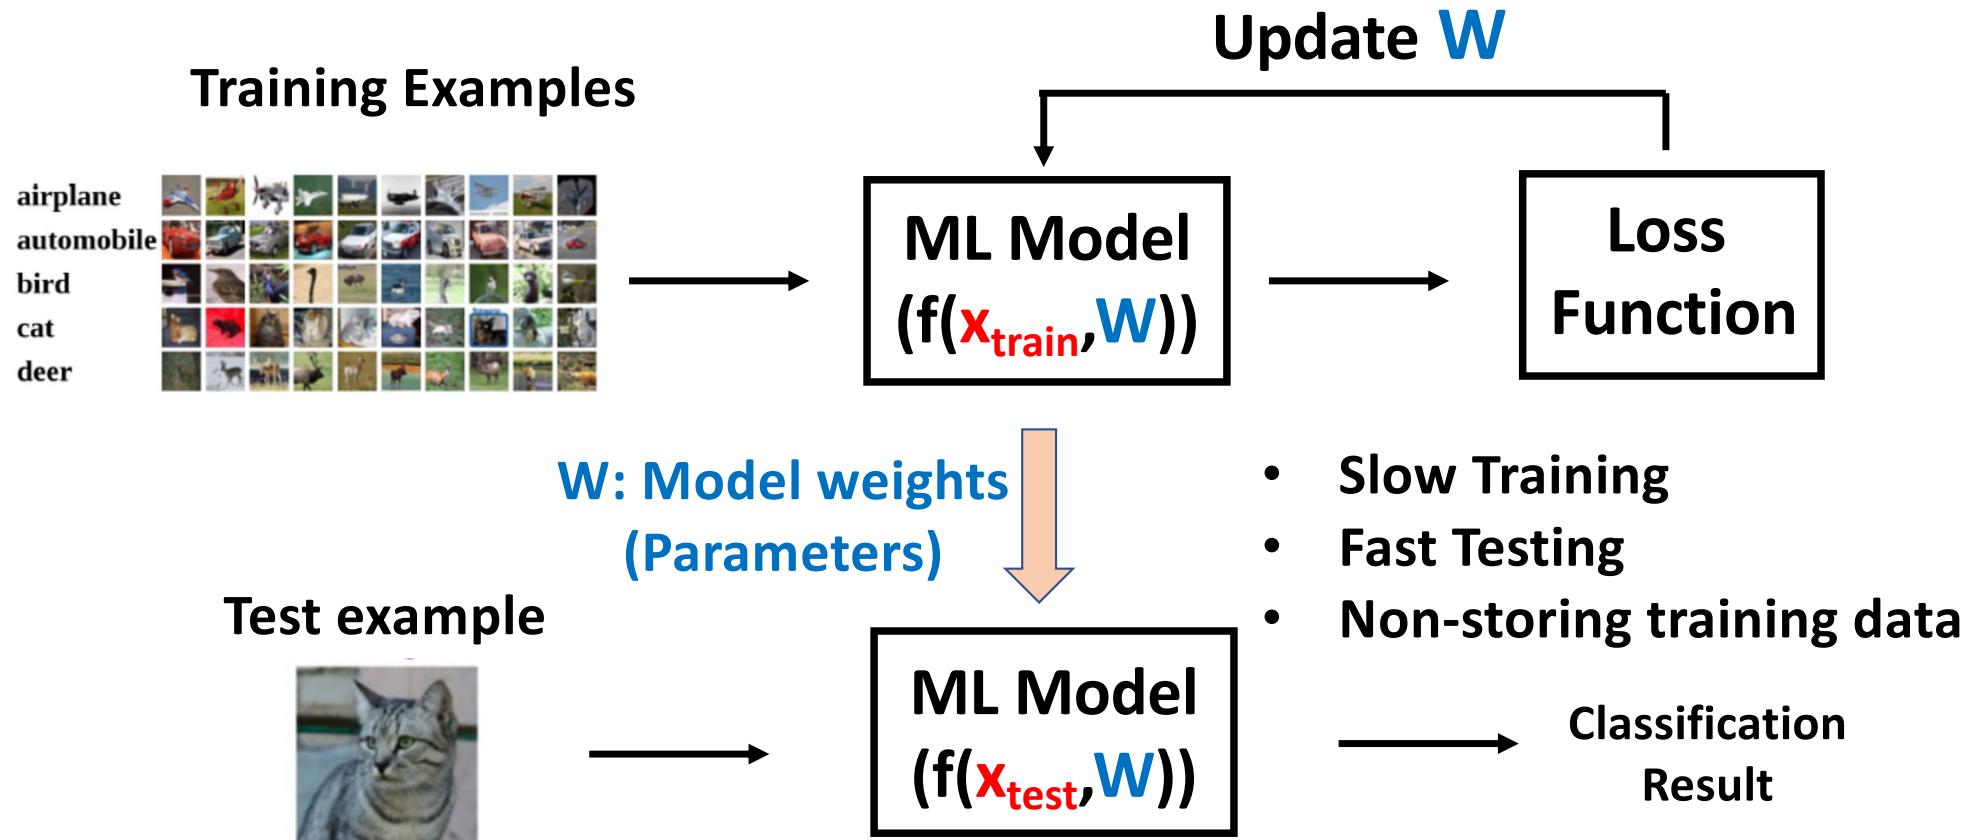
- **Pros:** Fast Training
Easy to implement
- **Cons:** Slow Testing
Store all training data

Test example

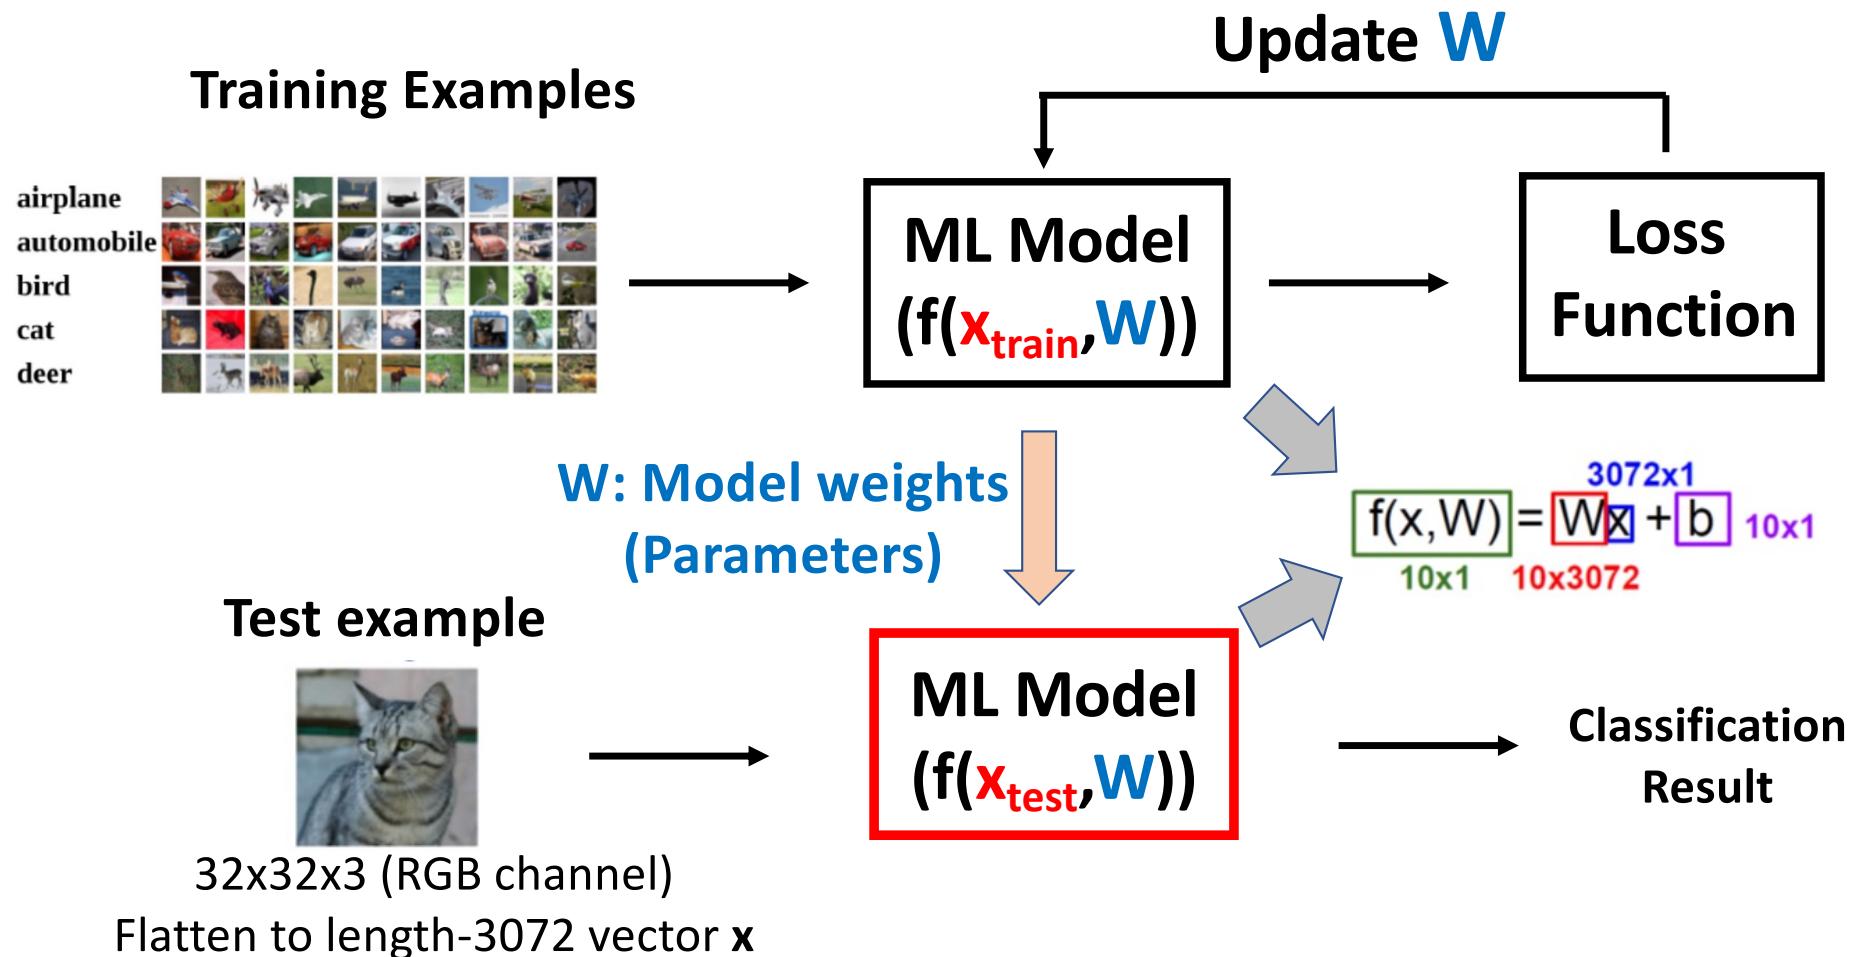


Classification
Result

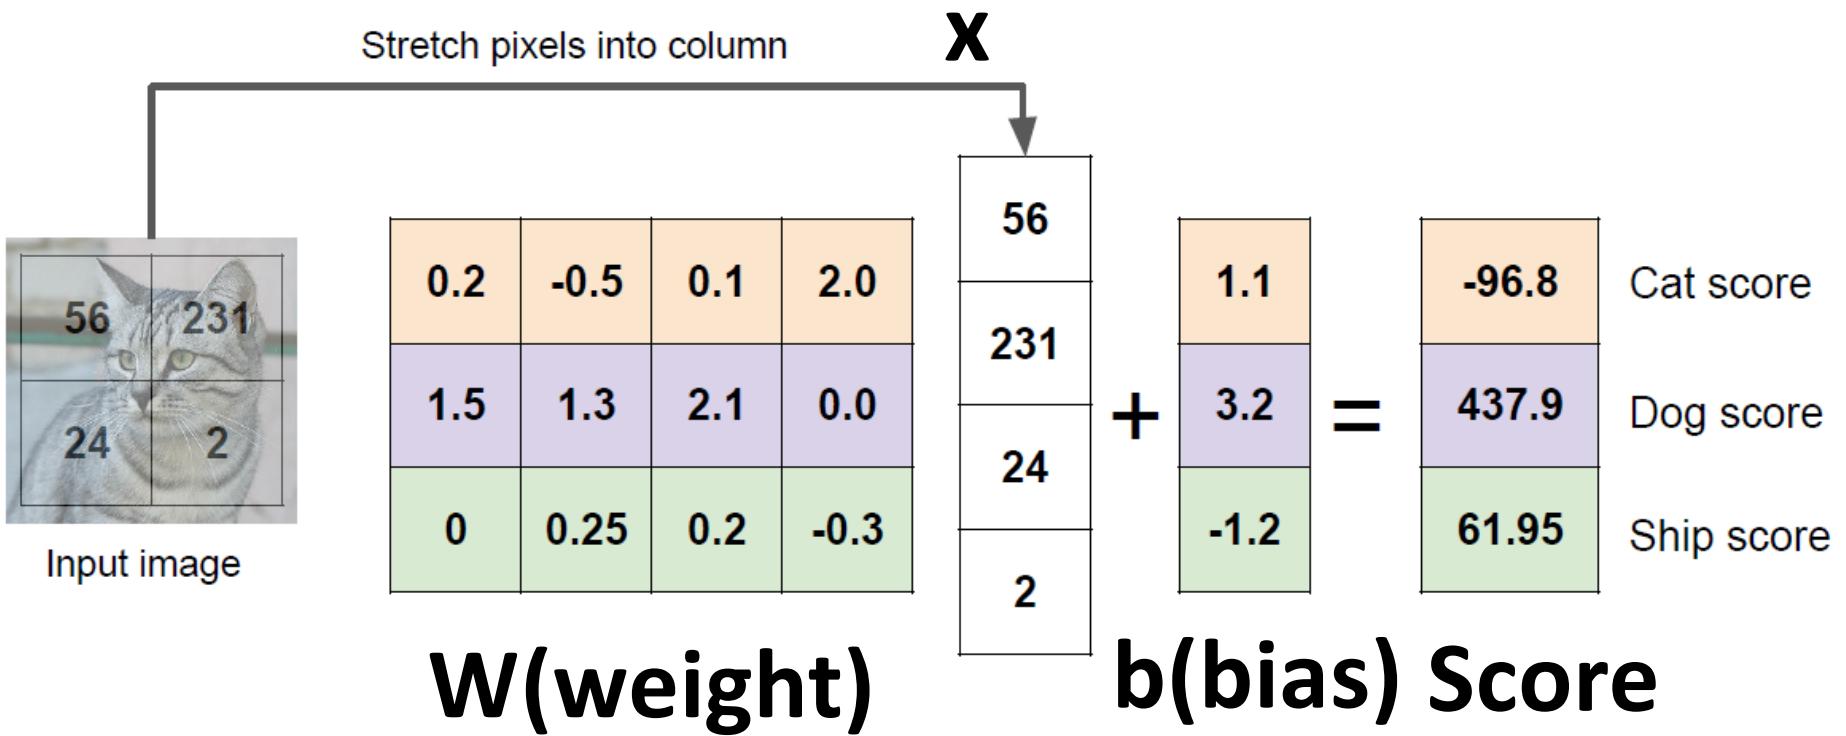
Parametric Machine Learning Model



Linear Classifier

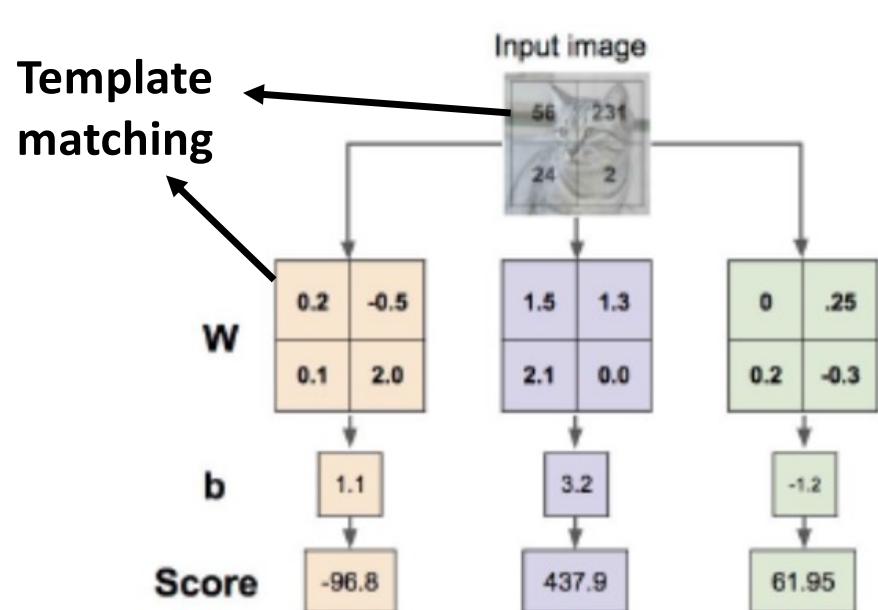


Vectorized Example



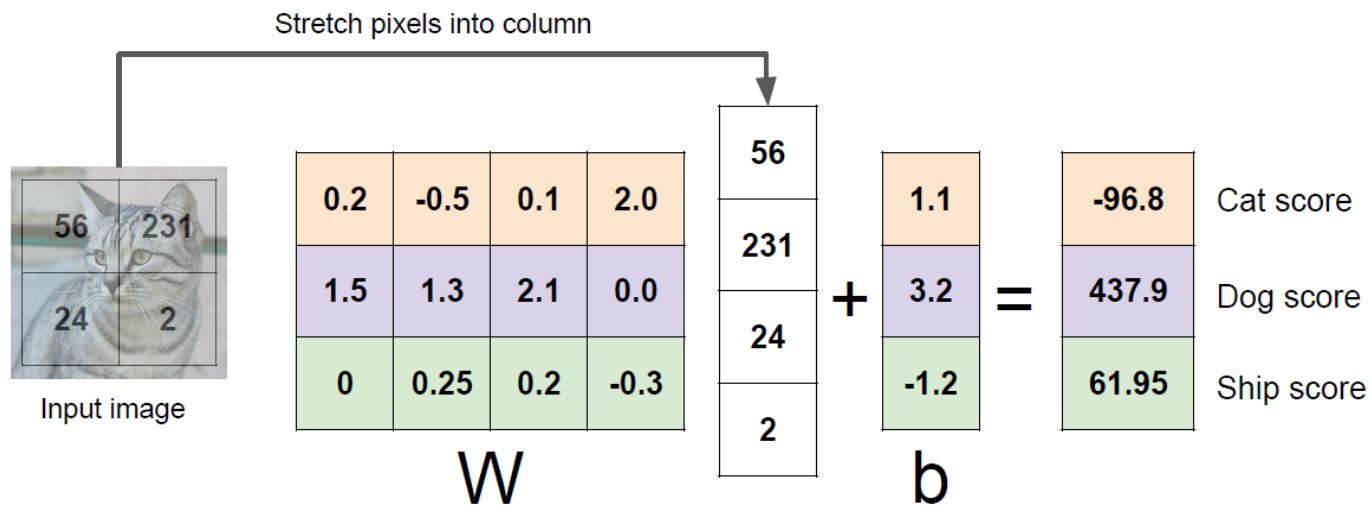
Q: Is it a well-trained model?

Visual Viewpoint of a Linear Classifier



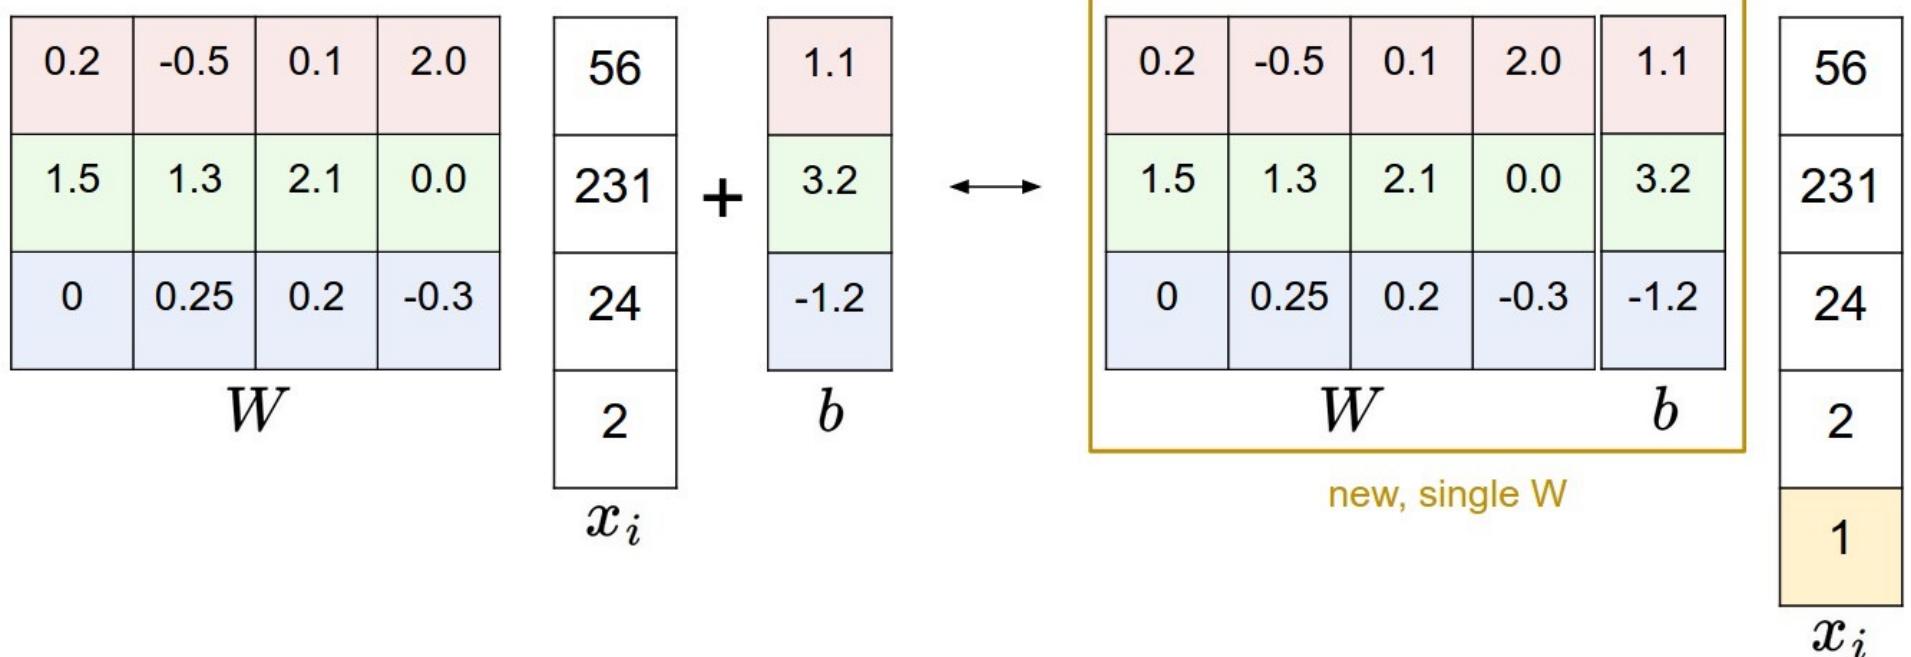
Why We Need Bias

- What happens if inputs are all zeros without bias?



- Viewpoint of function approximation
 - Linear classifier is to approximate an (assumed) linear relationship

Weights and Bias can be Unified



Acknowledgement

Many materials of the slides of this course are adopted and re-produced from several deep learning courses and tutorials.

- Prof. Fei-fei Li, Stanford, CS231n: Convolutional Neural Networks for Visual Recognition (online available)
- Prof. Andrew Ng, Stanford, CS230: Deep learning (online available)
- Prof. Yanzhi Wang, Northeastern, EECE7390: Advance in deep learning
- Prof. Jianting Zhang, CUNY, CSc G0815 High-Performance Machine Learning: Systems and Applications
- Prof. Vivienne Sze, MIT, “Tutorial on Hardware Architectures for Deep Neural Networks”
- Pytorch official tutorial <https://pytorch.org/tutorials/>