

14.332.435/16.332.530
Introduction to Deep Learning


Lecture 9
Training Neural Network 2

Yuqian Zhang

Department of Electrical and Computer Engineering

Babysitting Learning Process


- Step 1: Data Preprocessing



(Assume X [NxD] is data matrix,
each example in a row)

Babysitting Learning Process

- Step 2: Choose Architecture



Babysitting Learning Process

- Step 3a: Initial check the loss

```
def init_two_layer_model(input_size, hidden_size, output_size):  
    # initialize a model  
    model = {}  
    model['W1'] = 0.0001 * np.random.randn(input_size, hidden_size)  
    model['b1'] = np.zeros(hidden_size)  
    model['W2'] = 0.0001 * np.random.randn(hidden_size, output_size)  
    model['b2'] = np.zeros(output_size)  
    return model
```

```
model = init_two_layer_model(32*32*3, 50, 10) # input size, hidden size, number of classes  
loss, grad = two_layer_net(X_train, model, y_train, 0.0) # disable regularization  
print loss
```

2.30261216167

loss ~2.3.
"correct" for
10 classes

returns the loss and the
gradient for all parameters

Babysitting Learning Process

- Step 3b: Initial check the loss

```
def init_two_layer_model(input_size, hidden_size, output_size):  
    # initialize a model  
    model = {}  
    model['W1'] = 0.0001 * np.random.randn(input_size, hidden_size)  
    model['b1'] = np.zeros(hidden_size)  
    model['W2'] = 0.0001 * np.random.randn(hidden_size, output_size)  
    model['b2'] = np.zeros(output_size)  
    return model
```

```
model = init_two_layer_model(32*32*3, 50, 10) # input size, hidden size, number of classes  
loss, grad = two_layer_net(X_train, model, y_train, 1e3) # crank up regularization  
print loss
```

3.06859716482

loss went up, good. (sanity check)

Babysitting Learning Process

- Step 4: Use very small data to exam train process

```
model = init_two_layer_model(32*32*3, 50, 10) # input size, hidden size, number of classes
trainer = ClassifierTrainer()
```

```
X_tiny = X_train[:20] # take 20 examples
y_tiny = y_train[:20]
```

Use very small data

```
best_model, stats = trainer.train(X_tiny, y_tiny, X_tiny, y_tiny,
```

```
model, two_layer_net,
```

```
num_epochs=200, reg=0.0,
```

Turn off regularization

```
update='sgd', learning_rate_decay=1,
```

```
sample_batches = False,
```

```
learning_rate=1e-3, verbose=True)
```

Use simple optimizer

```
Finished epoch 1 / 200: cost 2.302603, train: 0.400000, val 0.400000, lr 1.000000e-03
Finished epoch 2 / 200: cost 2.302258, train: 0.450000, val 0.450000, lr 1.000000e-03
Finished epoch 3 / 200: cost 2.301849, train: 0.600000, val 0.600000, lr 1.000000e-03
Finished epoch 4 / 200: cost 2.301196, train: 0.650000, val 0.650000, lr 1.000000e-03
Finished epoch 5 / 200: cost 2.300044, train: 0.650000, val 0.650000, lr 1.000000e-03
Finished epoch 6 / 200: cost 2.297864, train: 0.550000, val 0.550000, lr 1.000000e-03
Finished epoch 7 / 200: cost 2.293595, train: 0.600000, val 0.600000, lr 1.000000e-03
Finished epoch 8 / 200: cost 2.285096, train: 0.550000, val 0.550000, lr 1.000000e-03
Finished epoch 9 / 200: cost 2.268094, train: 0.550000, val 0.550000, lr 1.000000e-03
Finished epoch 10 / 200: cost 2.234787, train: 0.500000, val 0.500000, lr 1.000000e-03
Finished epoch 11 / 200: cost 2.173187, train: 0.500000, val 0.500000, lr 1.000000e-03
Finished epoch 12 / 200: cost 2.076862, train: 0.500000, val 0.500000, lr 1.000000e-03
Finished epoch 13 / 200: cost 1.974098, train: 0.400000, val 0.400000, lr 1.000000e-03
Finished epoch 14 / 200: cost 1.895885, train: 0.400000, val 0.400000, lr 1.000000e-03
Finished epoch 15 / 200: cost 1.820876, train: 0.450000, val 0.450000, lr 1.000000e-03
Finished epoch 16 / 200: cost 1.737430, train: 0.450000, val 0.450000, lr 1.000000e-03
Finished epoch 17 / 200: cost 1.642356, train: 0.500000, val 0.500000, lr 1.000000e-03
Finished epoch 18 / 200: cost 1.535239, train: 0.600000, val 0.600000, lr 1.000000e-03
Finished epoch 19 / 200: cost 1.421527, train: 0.600000, val 0.600000, lr 1.000000e-03
```

```
Finished epoch 195 / 200: cost 0.002694, train: 1.000000, val 1.000000, lr 1.000000e-03
Finished epoch 196 / 200: cost 0.002674, train: 1.000000, val 1.000000, lr 1.000000e-03
Finished epoch 197 / 200: cost 0.002655, train: 1.000000, val 1.000000, lr 1.000000e-03
Finished epoch 198 / 200: cost 0.002635, train: 1.000000, val 1.000000, lr 1.000000e-03
Finished epoch 199 / 200: cost 0.002617, train: 1.000000, val 1.000000, lr 1.000000e-03
```

Make sure for very small portion we can overfit

Babysitting Learning Process

- Step 4: Begin to train via trying small regularization and learning rate to reduce loss

```
model = init_two_layer_model(32*32*3, 50, 10) # input size, hidden size, number of classes
trainer = ClassifierTrainer()
best_model, stats = trainer.train(X_train, y_train, X_val, y_val,
                                  model, two_layer_net,
                                  num_epochs=10, reg=0.000001,
                                  update='sgd', learning_rate_decay=1,
                                  sample_batches = True,
                                  learning_rate=1e-6, verbose=True)
```

```
Finished epoch 1 / 10: cost 2.302576, train: 0.080000, val 0.103000, lr 1.000000e-06
Finished epoch 2 / 10: cost 2.302582, train: 0.121000, val 0.124000, lr 1.000000e-06
Finished epoch 3 / 10: cost 2.302558, train: 0.119000, val 0.138000, lr 1.000000e-06
Finished epoch 4 / 10: cost 2.302519, train: 0.127000, val 0.151000, lr 1.000000e-06
Finished epoch 5 / 10: cost 2.302517, train: 0.158000, val 0.171000, lr 1.000000e-06
Finished epoch 6 / 10: cost 2.302518, train: 0.179000, val 0.172000, lr 1.000000e-06
Finished epoch 7 / 10: cost 2.302466, train: 0.180000, val 0.176000, lr 1.000000e-06
Finished epoch 8 / 10: cost 2.302452, train: 0.175000, val 0.185000, lr 1.000000e-06
Finished epoch 9 / 10: cost 2.302459, train: 0.206000, val 0.192000, lr 1.000000e-06
Finished epoch 10 / 10: cost 2.302420, train: 0.190000, val 0.192000, lr 1.000000e-06
finished optimization. best validation accuracy: 0.192000
```

**Too small learning
rate makes loss
barely down**

Babysitting Learning Process

- Step 4: Begin to train via trying small regularization and learning rate to reduce loss

```
model = init_two_layer_model(32*32*3, 50, 10) # input size, hidden size, number of classes
trainer = ClassifierTrainer()
best_model, stats = trainer.train(X_train, y_train, X_val, y_val,
                                  model, two_layer_net,
                                  num_epochs=10, reg=0.000001,
                                  update='sgd', learning_rate_decay=1,
                                  sample_batches = True,
                                  learning_rate=1e5, verbose=True)
```



/home/karpathy/cs231n/code/cs231n/classifiers/neural_net.py:50: RuntimeWarning: divide by zero encountered in log
data_loss = -np.sum(np.log(probs[range(N), y])) / N
/home/karpathy/cs231n/code/cs231n/classifiers/neural_net.py:48: RuntimeWarning: invalid value encountered in subtract
probs = np.exp(scores - np.max(scores, axis=1, keepdims=True))

Finished epoch 1 / 10: cost nan, train: 0.091000, val 0.087000, lr 1.000000e+06
Finished epoch 2 / 10: cost nan, train: 0.095000, val 0.087000, lr 1.000000e+06
Finished epoch 3 / 10: cost nan, train: 0.100000, val 0.087000, lr 1.000000e+06

Too high learning rate makes loss explode



Proper Step Size is Important


- Too small: Long training time
- Too big: Skip optimal point (hard to converge)




Notes on Training

- Typical rang of learning rate
 - within the range $[1e-3 \dots 1e-5]$
- How to determine the end of training
 - set maximum iteration
 - observe convergence of loss
 - observe validation error rate





Overfit and Underfit



big gap = overfitting
=> increase regularization strength?

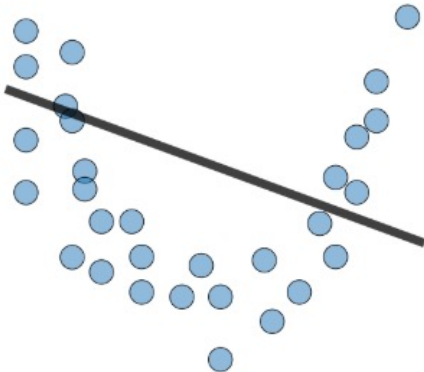
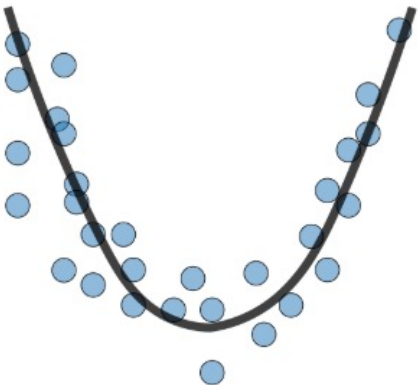
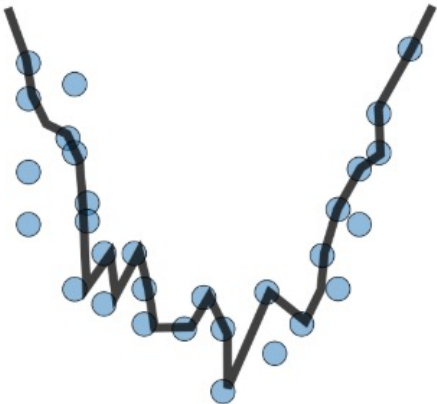
no gap
=> increase model capacity?

Bias and Variance

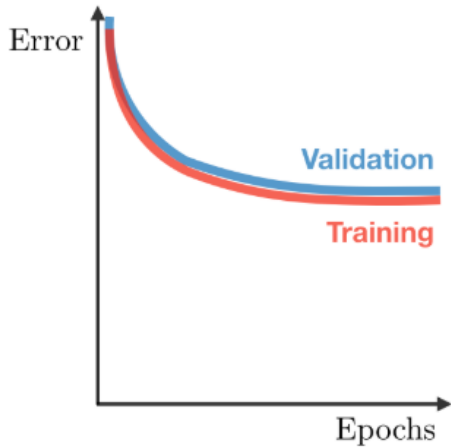
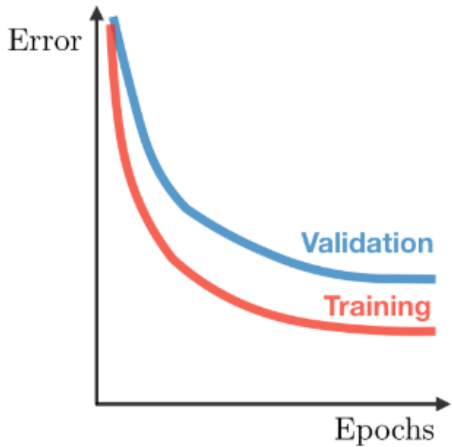
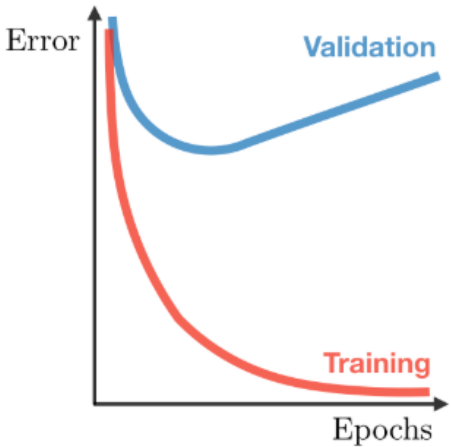
- The *bias* is an error from erroneous assumptions in the learning *algorithm*. High bias can cause an algorithm to miss the relevant relations between features and target outputs (underfitting).
- The *variance* is an error from sensitivity to small fluctuations in the training set. High variance can cause an algorithm to model the random *noise* in the training data, rather than the intended outputs (*overfitting*).

In *statistics* and *machine learning*, the **bias–variance tradeoff** is the property of a set of predictive models whereby models with a lower *bias* in *parameter estimation* have a higher *variance* of the parameter estimates across *samples*, and vice versa. The **bias–variance**

From wiki



	Underfitting	Just right	Overfitting
Symptoms	<ul style="list-style-type: none"> • High training error • Training error close to test error • High bias 	<ul style="list-style-type: none"> • Training error slightly lower than test error 	<ul style="list-style-type: none"> • Very low training error • Training error much lower than test error • High variance
Regression illustration			

	Underfitting	Just right	Overfitting
Symptoms	<ul style="list-style-type: none"> • High training error • Training error close to test error • High bias 	<ul style="list-style-type: none"> • Training error slightly lower than test error 	<ul style="list-style-type: none"> • Very low training error • Training error much lower than test error • High variance
Classification illustration			


	Underfitting	Just right	Overfitting
Symptoms	<ul style="list-style-type: none"> • High training error • Training error close to test error • High bias 	<ul style="list-style-type: none"> • Training error slightly lower than test error 	<ul style="list-style-type: none"> • Very low training error • Training error much lower than test error • High variance
Deep learning illustration	 <p>The graph shows Error on the y-axis and Epochs on the x-axis. A red line (Training) and a blue line (Validation) both start at a high error level and decrease very slowly over time, remaining close to each other. This indicates the model is not learning effectively from the data.</p>	 <p>The graph shows Error on the y-axis and Epochs on the x-axis. A red line (Training) starts high and decreases smoothly to a low level. A blue line (Validation) starts slightly higher than the training error and decreases to a level just above the training error, indicating a well-balanced model.</p>	 <p>The graph shows Error on the y-axis and Epochs on the x-axis. A red line (Training) starts high and decreases rapidly to a very low level. A blue line (Validation) starts high, decreases to a minimum, and then begins to rise steadily, indicating the model is memorizing the training data and failing to generalize.</p>
Possible remedies	<ul style="list-style-type: none"> • Complexify model • Add more features • Train longer 		<ul style="list-style-type: none"> • Perform regularization • Get more data

<https://stanford.edu/~shervine/teaching/cs-229/cheatsheet-machine-learning-tips-and-tricks#>

Early Stopping



Problem of SGD



Local Minima

Saddle Point

Zero gradient, gradient descent gets stuck

SGD with Momentum

Think an analogy

- Interpret loss as the height (h) of a hilly terrain
- Interpret weight initialization as setting a particle with zero initial velocity (v) at some location.
- Interpret optimization process as simulating the particle rolling on the landscape.
- Interpret the gradient as the force felt by the particle
 - Because $U = mgh$ $F = -\nabla U$

SGD with Momentum

Now consider $F = ma$ $U = mgh$ $F = -\nabla U$

- Gradient can be viewed proportional the acceleration (a) of the particle.
- Conventional SGD gradient directly integrates the position.
- This analogy implies gradient only directly influences the velocity, which in turn has an effect on the position

SGD with Momentum

Conventional SGD

```
# Vanilla update  
x += - learning_rate * dx
```

SGD with Momentum

```
# Momentum update  
v = mu * v - learning_rate * dx # integrate velocity  
x += v # integrate position
```

v initialized as 0

mu: momentum, hyperparameter, usually 0.9

After using SGD+Momentum

Local Minima




Saddle points



Problem of SGD


- When loss changes quickly in one direction and slowly in another
 - Very slow progress along shallow dimension, jitter along steep direction



Adjust Learning Rate Adaptively

- Avoid globally equal learning rate can alleviate

$$w = w - \textit{lamda} * dw$$



AdaGrad

- Element-wise adaptively adjust *effective* learning rate
 - for weights that receive high gradients, reduce
 - for weights that receive small or infrequent update, increase

```
# Assume the gradient dx and parameter vector x  
cache += dx**2  
x += - learning_rate * dx / (np.sqrt(cache) + eps)
```

eps is between 1e-4 and 1e-8

RMSprop (Root Mean Squared Propagation)

- Cons of AdaGrad
 - Learning rate monotonically decreases
 - Stops learning too early
- RMSprop avoid this via using moving average

```
cache = decay_rate * cache + (1 - decay_rate) * dx**2
x += - learning_rate * dx / (np.sqrt(cache) + eps)
```

decay rate is typically 0.9, 0.99, 0.999...

Note about Learning Rate

- (Global) Learning rate should also be decayed gradually

step decay:


e.g. decay learning rate by half every few epochs.

exponential decay:

$$\alpha = \alpha_0 e^{-kt}$$

1/t decay:

$$\alpha = \alpha_0 / (1 + kt)$$




Pytorch Learning Rate Scheduler

Decreases the **learning rate** by **gamma** every **step_size** epochs.

- `from torch.optim.lr_scheduler import StepLR`
- `optimizer = torch.optim.SGD(model.parameters(), lr=0.1)`
- `scheduler = StepLR(optimizer, step_size=30, gamma=0.1)`
- `for epoch in range(num_epochs):`
 - `train(...)`
 - `scheduler.step()`


First-Order Optimization

- Gradient descent belongs to first-order optimization
- Convergence rate is relatively slow



Second-Order Optimization

- Second-order optimization has fast convergence rate



Pros of Second-Order Optimization

second-order Taylor expansion:

Hessian matrix

$$J(\boldsymbol{\theta}) \approx J(\boldsymbol{\theta}_0) + (\boldsymbol{\theta} - \boldsymbol{\theta}_0)^\top \nabla_{\boldsymbol{\theta}} J(\boldsymbol{\theta}_0) + \frac{1}{2} (\boldsymbol{\theta} - \boldsymbol{\theta}_0)^\top \mathbf{H} (\boldsymbol{\theta} - \boldsymbol{\theta}_0)$$

Solving for the critical point we obtain the Newton parameter update:

$$\boldsymbol{\theta}^* = \boldsymbol{\theta}_0 - \mathbf{H}^{-1} \nabla_{\boldsymbol{\theta}} J(\boldsymbol{\theta}_0)$$

No hyperparameters!

No learning rate!

(Though you might use one in practice)

Suppose $f: \mathbb{R}^n \rightarrow \mathbb{R}$ is a function taking as input a vector $\mathbf{x} \in \mathbb{R}^n$ and outputting a scalar $f(\mathbf{x}) \in \mathbb{R}$; if all second [partial derivatives](#) of f exist and are continuous over the domain of the function, then the Hessian matrix \mathbf{H} of f is a square $n \times n$ matrix, usually defined and arranged as follows:

$$\mathbf{H} = \begin{bmatrix} \frac{\partial^2 f}{\partial x_1^2} & \frac{\partial^2 f}{\partial x_1 \partial x_2} & \cdots & \frac{\partial^2 f}{\partial x_1 \partial x_n} \\ \frac{\partial^2 f}{\partial x_2 \partial x_1} & \frac{\partial^2 f}{\partial x_2^2} & \cdots & \frac{\partial^2 f}{\partial x_2 \partial x_n} \\ \vdots & \vdots & \ddots & \vdots \\ \frac{\partial^2 f}{\partial x_n \partial x_1} & \frac{\partial^2 f}{\partial x_n \partial x_2} & \cdots & \frac{\partial^2 f}{\partial x_n^2} \end{bmatrix}.$$

or, by stating an equation for the coefficients using indices i and j :

$$\mathbf{H}_{i,j} = \frac{\partial^2 f}{\partial x_i \partial x_j}.$$

From wiki

Cons of Second-Order Optimization

second-order Taylor expansion:

$$J(\boldsymbol{\theta}) \approx J(\boldsymbol{\theta}_0) + (\boldsymbol{\theta} - \boldsymbol{\theta}_0)^\top \nabla_{\boldsymbol{\theta}} J(\boldsymbol{\theta}_0) + \frac{1}{2} (\boldsymbol{\theta} - \boldsymbol{\theta}_0)^\top \mathbf{H} (\boldsymbol{\theta} - \boldsymbol{\theta}_0)$$

Solving for the critical point we obtain the Newton parameter update:

$$\boldsymbol{\theta}^* = \boldsymbol{\theta}_0 - \mathbf{H}^{-1} \nabla_{\boldsymbol{\theta}} J(\boldsymbol{\theta}_0)$$

Hessian has $O(N^2)$ elements

Inverting takes $O(N^3)$

N = (Tens or Hundreds of) Millions

Acknowledgement

Many materials of the slides of this course are adopted and re-produced from several deep learning courses and tutorials.

- Prof. Fei-fei Li, Stanford, CS231n: Convolutional Neural Networks for Visual Recognition (online available)
- Prof. Andrew Ng, Stanford, CS230: Deep learning (online available)
- Prof. Yanzhi Wang, Northeastern, EECE7390: Advance in deep learning
- Prof. Jianting Zhang, CUNY, CSc G0815 High-Performance Machine Learning: Systems and Applications
- Prof. Vivienne Sze, MIT, “Tutorial on Hardware Architectures for Deep Neural Networks”
- Pytorch official tutorial <https://pytorch.org/tutorials/>
- <https://github.com/jcjohnson/pytorch-examples>